
Research Report: Current primary-source comparison for safe multi-agent orchestration of a production incident response system: OpenAI Agents SDK, CrewAI, LangGraph, Microsoft Agent Framework or AutoGen, Google ADK, Temporal durable workflows. Focus on supervisor/router, handoffs, event-driven ledger/blackboard, graph state machines, durable execution, human approval, evaluator isolation, testing/backtesting, retries and stop conditions. Return official documentation URLs and distinguish facts from recommendations.

- Pack: generic
- Profile: generic
- As of: 2026-07-17
- Run status: failed_no_rows
- Raw rows: 0
- Stance: no_evidence_collected
- Confidence: low
- Action bias: analyze_before_action
- Average evidence quality: 0.0
- Duplicate clusters: 0
- Conflict flags: 0

Executive Summary

Collected 0 eligible row(s) from 0 total across 0 connector(s).

- 0 claim buckets are supported.
- Matrix gap assessment: not_applicable.

Scope and Coverage

- Required facets: 5

- Required facets covered: 0
- Missing required facets: overview, primary_sources, current_evidence, alternatives, risks

Evidence

Evidence Quality Warnings

- web_search anysearch HTTP 402

Loop Status

- Loop status: blocked
- Stop reason: sources_returned_no_evidence
- Feedback actions:
 - failed_no_rows: Broaden --platform-scope, simplify the query, or repair missing/auth-expired connectors.
 - stop_brakes: Set --source-timeout-seconds so long-running connectors have a hard stop.
 - bounded_execution: Set --source-timeout-seconds for fragile or long-running connectors.
 - evidence_collected: Broaden collection or repair connectors.
 - source_quality: Collect evidence before synthesis.